

Class 2: Regression and classification

Andrew Parnell, School of Mathematics and Statistics,
University College Dublin

Learning outcomes

- ▶ Be able to understand the structure of regression and classification models
- ▶ Know how to read and interpret the output of a statistical model
- ▶ Be familiar with some of the extensions to basic regression and classification models

Why regression and classification?

- ▶ t-tests are only really useful when you have a continuous outcome variable and one discrete variable with two groups (e.g. treatment vs control)
- ▶ For almost any real life situation you have multiple variables of all different types
- ▶ For these situations you need a *statistical model*
- ▶ A statistical model allows us to perform *probabilistic prediction* of the outcome variable from the remaining variables, and/or to explain how the other variables are causing the outcome variable to change

Regression vs Classification: what's the difference?


- ▶ In regression we have a single *continuous* outcome variable and lots of other variables which we think might be good predictors of the outcome
- ▶ In classification we have a single *discrete* outcome variable and lots of other variables
- ▶ In the machine learning literature this is often known as *supervised learning*
- ▶ Situations where there are multiple outcome variables are beyond the scope of this course

Response and explanatory variables

- ▶ The outcome variable is more commonly known as the *response* variable
- ▶ The other variables which we think might be good predictors of the response variable are called the *explanatory variables* (though be careful with causation)
- ▶ We will use these words from now on, but beware there are lots of other terms in the literature

A basic regression model

- ▶ Let's go back to the prostate cancer data
- ▶ Recall the key outcome variable is `lpsa` the log of the prostate specific antigen value. This is our response variable
- ▶ Suppose we had one explanatory variable `lweight`



Creating the model


- ▶ Looking at the plot, there may be a positive, linear relationship between $\log(\text{weight})$ and $\log(\text{psa})$
- ▶ Perhaps we can create a prediction model that allows us to predict $\log(\text{psa})$ from $\log(\text{weight})$
- ▶ Suppose, for each patient we multiplied the $\log(\text{weight})$ value by 1.2 and then subtracted the value 2 so:

$$\text{prediction} = 1.2 \times \log(\text{weight}) - 2$$

- ▶ If we do this repeatedly for every value in the data set we get
...

A first model

```
prediction = 1.2 * prostate$weight - 2  
plot(prostate$weight, prostate$psa,  
      xlab = 'log(weight)', ylab = 'log(psa)')  
points(prostate$weight, prediction, col='red')
```



Refining the model

- ▶ Is this model any good?
- ▶ How might we measure how close our predictions are to the truth?
- ▶ How can we choose the values (here 1.2 and -2) better?

Getting R to do the work

- Luckily the R function `lm` will do the work for us

```
model_1 = lm(formula = lpsa ~ lweight, data = prostate)
summary(model_1)
```

```
...
coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  -1.7586      0.9103  -1.932   0.0564 .
lweight       1.1676      0.2491   4.686 9.28e-06 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.046 on 95 degrees of freedom
Multiple R-squared:  0.1878, Adjusted R-squared:  0.1792
F-statistic: 21.96 on 1 and 95 DF,  p-value: 9.276e-06
...
```


Background details

- ▶ The two values here are the y -intercept and the slope of the line. They are commonly known as the *regression coefficients*
- ▶ R chooses these coefficients by minimising the vertical distances between the black and the red points
- ▶ A key assumption in the model is that these vertical distances (known as *residuals*) are normally distributed
- ▶ R uses this assumption to run t-tests on the parameters, which you can see the results of in the `summary` output

Plotting the fit

One way is to type `plot(model_1)` which gives residual diagnostics. A quick plot of the fitted line via:

```
plot(prostate$lweight, prostate$lpsa,  
     xlab = 'log(weight)', ylab= 'log(psa)')  
abline(model_1, col='red')
```



Expanding the model with two explanatory variables

Suppose we wanted to use two explanatory variables, lweight and age:

```
model_2 = lm(formula = lpsa ~ lweight + age,  
              data = prostate)  
summary(model_2)
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-1.897709	1.119033	-1.696	0.0932	.
lweight	1.147487	0.267094	4.296	4.23e-05	***
age	0.003318	0.015369	0.216	0.8295	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.051 on 94 degrees of freedom

Multiple R-squared: 0.1882, Adjusted R-squared: 0.1709

F-statistic: 10.89 on 2 and 94 DF, p-value: 5.558e-05

Expanding the fit even more

```
model_3 = lm(formula = lpsa ~ . - train, data = prostate)
summary(model_3)
```

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	0.181561	1.320568	0.137	0.89096	
lcavol	0.564341	0.087833	6.425	6.55e-09	***
lweight	0.622020	0.200897	3.096	0.00263	**
age	-0.021248	0.011084	-1.917	0.05848	.
lbph	0.096713	0.057913	1.670	0.09848	.
svi	0.761673	0.241176	3.158	0.00218	**
lcp	-0.106051	0.089868	-1.180	0.24115	
gleason	0.049228	0.155341	0.317	0.75207	
pgg45	0.004458	0.004365	1.021	0.31000	

Multiple regression

- ▶ When you have lots of explanatory variables this is known as *multiple regression*
- ▶ You can still use the values in the Estimate column to create predictions of lpsa by multiplying and adding up
- ▶ Beware the p-values as before: they might be highly significant but still a very poor model
- ▶ R gives you two other useful statistics:
 - ▶ The R-squared which measures the proportion of variation in the response variable explained by the explanatory variables
 - ▶ The residual standard error which measures how far away the data points are from the fitted line

Dealing with interactions

- ▶ Our explanatory variables will often interact with each other to affect the response variable
- ▶ The usual way to deal with interactions is to create *new explanatory variables* by multiplying them together. `lm` does this for you:

```
model_4 = lm(formula = lpsa ~ lweight + age + lweight:age,  
summary(model_4)
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-9.97325	8.45553	-1.179	0.241
lweight	3.45163	2.40620	1.434	0.155
age	0.12575	0.12800	0.982	0.328
lweight:age	-0.03481	0.03613	-0.964	0.338

Residual standard error: 1.051 on 93 degrees of freedom
Multiple R-squared: 0.1962, Adjusted R-squared: 0.1703
F-statistic: 7.566 on 3 and 93 DF, p-value: 0.0001391

Final remarks on regression models

- ▶ There is lots of research on regression models of all different types
- ▶ The vast majority of them involve creating a set of coefficients to multiply the explanatory variables by and then adding everything up
- ▶ It becomes very hard to plot the predictions in large and complex models
- ▶ It's very important to check the model diagnostics using `plot` and to look at the R-squared and residual standard error values

Classification models

Intro to classification models

- ▶ Returning to the South African heart rate data, recall that here we are interested in predicting whether someone has CHD or not
- ▶ We have explanatory variables including adiposity, alcohol use, age, etc
- ▶ CHD is a discrete binary variable (1 or 0). It thus makes more sense to try and predict a probability of CHD i.e. a value between 0 and 1, rather than CHD itself
- ▶ If we use our previous approach to guess coefficients for the different explanatory variables we will run into problems with values going outside 1 or 0

The logit transformation

- ▶ Suppose we wanted to predict CHD from age
- ▶ We might come up with the model:

$$Prob(CHD) = 0.06 \times age - 2$$

- ▶ Thus if someone has an age of 40 they have probability 0.4
- ▶ But if someone has an age of 20 they have probability -0.8. Oh dear!
- ▶ Instead use a transformation, such as the *logit*

$$Prob(CHD) = \frac{\exp(0.06 \times age - 2)}{1 + \exp(0.06 \times age - 2)}$$

- ▶ This transformation *garuantees* the values will be between 0 and 1 - try it!

About classification models

- ▶ Rather than try to predict a continuous response variable, classification models aim to find the probability that an observation is in a particular class
- ▶ Underneath the hood though, they are exactly like regression models with coefficients applied to each of the explanatory variables before adding everything up
- ▶ We then use a clever transformation (such as the logit, but there are others) to turn it into a probability
- ▶ Rather than the normal distribution, we use the *binomial* distribution to judge how close the observations are to the predictions and hence estimate the missing coefficients
- ▶ The R function `glm` will fit a classification model for us

Example: SA Heart rate

```
model_1 = glm(chd ~ age, data = SA, family = 'binomial')
summary(model_1)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	-3.521710	0.416031	-8.465	< 2e-16	***
age	0.064108	0.008532	7.513	5.76e-14	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 596.11 on 461 degrees of freedom
Residual deviance: 525.56 on 460 degrees of freedom
AIC: 529.56

Understanding the output

- ▶ The output here is much less helpful
- ▶ We have the coefficient values, but this is before the logit transformation so not particularly useful
- ▶ We have the p-values of the coefficients but we should be wary of these
- ▶ The other values (deviance etc) aren't particularly helpful
- ▶ AIC we'll talk about in the next class
- ▶ In fact, to judge the performance of the model we need to do a lot more work!

Extending the model

We can extend to multiple explanatory variables in exactly the same way as before:


```
model_2 = glm(chd ~ age + adiposity + age:adiposity, data =  
summary(model_2)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	-4.6909894	1.3851412	-3.387	0.000708	***
age	0.0811012	0.0300150	2.702	0.006892	**
adiposity	0.0583492	0.0596028	0.979	0.327596	
age:adiposity	-0.0009184	0.0012051	-0.762	0.446000	

Plotting the fitted model

```
plot(jitter(SA$age), SA$chd, ylab = 'chd', xlab = 'Age')  
points(SA$age, model_1$fitted.values, col='red')
```



Regularisation and shrinkage

- ▶ When you have lots and lots of explanatory variables, the model can become very slow or might not fit at all
- ▶ Worse, we might have lots of spurious small p-values without any predictive power
- ▶ It makes sense to remove or reduce some of the coefficients on the explanatory variables if we think their effect is over-stated
- ▶ One way of doing this is via *regularisation*, where we set some of the values to zero, another is via *shrinkage* where we reduce the values (shrink them towards zero)

Lasso and Ridge

- ▶ The R package `glmnet` will perform shrinkage and regularisation for both regression and classification
- ▶ The *Lasso* model imposes a restricted sum on the absolute value of all of the coefficient values
- ▶ The *Ridge* model imposes an assumption that all of the coefficient values come from a normal distribution with some small standard deviation
- ▶ Fitting these types of model is beyond the scope of this course

More advanced classification approaches

- ▶ Much like regression, classification models have a long literature
- ▶ However, classification models tend to be more complicated as there are transformations involved (e.g. logit) and often multiple response variables (i.e. more than two categories for the response)
- ▶ Sometimes you have the choice between using a discrete response variable or a continuous one. I would always pick the continuous one: in general *regression models perform better than classification models*

Summary

- ▶ We now know how to implement regression and classification models in R
- ▶ We know how to interpret the output of some regression models
- ▶ We're familiar with some of the more advanced concepts in regression and classification